

## Distributed Database Systems

### ◆ Definition:

A **Distributed Database System (DDBS)** is a collection of **logically interrelated databases** distributed across multiple physical locations that appear to the user as a **single unified system**.

---

### ◆ 1. Key Features of Distributed Databases:

- **Data Distribution:** Data is stored across multiple sites or nodes.
  - **Location Transparency:** Users are unaware of the physical location of data.
  - **Concurrency Control:** Ensures correctness during simultaneous access.
  - **Fault Tolerance:** Handles node failures without loss of service.
  - **Autonomy:** Each site can manage its own database independently.
- 

### ◆ 2. Architectures of Distributed Database Systems

#### ✳ A. Client-Server Architecture

- **Client** sends queries.
- **Server** processes and returns data.
- Simple, centralized control over distribution.

#### ✳ B. Peer-to-Peer Architecture

- Each node is both a client and a server.
- Nodes collaborate for query execution.
- Better fault tolerance and scalability.

#### ✳ C. Multi-Database Architecture (Federated)

- Each database is managed independently.
- Integration occurs at a higher layer.
- Suitable when databases differ in types or platforms.

#### ✳ D. Cloud-Based DDBS

- Hosted on cloud infrastructure (e.g., AWS, Azure, GCP).
  - Elastic, globally distributed, pay-per-use.
  - Examples: **Google Spanner, Amazon Aurora, CockroachDB**.
- 

### ◆ 3. Types of Data Distribution:

Type	Description
<b>Fragmentation</b>	Dividing a table into smaller pieces (horizontally/vertically).
<b>Replication</b>	Copying data across multiple sites to improve availability.
<b>Allocation</b>	Determining where to place fragments and replicas for performance.

---

◆ **4. Consistency Models in Distributed Databases**

Consistency models define how updates to data are seen by different users and systems.

Model	Description
<b>Strong Consistency</b>	All nodes see the same data at the same time.
<b>Eventual Consistency</b>	Updates will propagate over time, and all nodes will eventually see the same data.
<b>Causal Consistency</b>	Writes that are causally related must be seen in order.
<b>Session Consistency</b>	Guarantees consistency within a single user session.
<b>Read-Your-Writes</b>	Ensures a user sees the results of their own writes.
<b>Quorum-Based Consistency</b>	Requires a majority of nodes to agree on a value (used in systems like DynamoDB).

---

◆ **5. Advantages of Distributed Databases:**

- Improved **reliability** and **availability**
  - Better **performance** via local access
  - **Scalability** across regions or data centers
  - **Modular growth** – nodes can be added easily
- 

◆ **6. Challenges:**

- **Network latency** and partitioning
  - **Distributed transaction management**
  - Maintaining **consistency** across replicas
  - **Security** across distributed infrastructure
- 

◆ **7. Real-World Examples:**

System	Type	Used By
<b>Google Spanner</b>	Globally distributed SQL	Google, Cloud customers
<b>Amazon DynamoDB</b>	NoSQL, eventual consistency	AWS services, real-time apps
<b>Apache Cassandra</b>	Peer-to-peer, NoSQL	Netflix, Instagram
<b>MongoDB Atlas</b>	Document-based	Web and mobile apps

### CAP Theorem (Brewer's Theorem)

- ◆ **Definition:**

The **CAP theorem** states that in any **distributed data system**, it is **impossible to simultaneously guarantee all three** of the following properties:

CAP Component	Explanation
<b>Consistency (C)</b>	Every read receives the most recent write or an error.
<b>Availability (A)</b>	Every request receives a (non-error) response, without guarantee of recent data.
<b>Partition Tolerance (P)</b>	The system continues to function even when communication between nodes is lost.

- ◆ **Implication:**

You can **only achieve two of the three** at any time:

System Type Guarantees	Example
CP	Consistency + Partition Tolerance HBase, MongoDB (in some configs)
CA	Consistency + Availability Rare in distributed systems
AP	Availability + Partition Tolerance DynamoDB, Couchbase, Cassandra

### 2. ACID Properties (Relational Databases)

- ◆ **Definition:**

ACID stands for a set of **guarantees** that traditional databases provide to ensure data integrity during transactions.

Property	Meaning
----------	---------

<b>Atomicity</b>	All operations in a transaction are completed; if one fails, all are rolled back.
------------------	---

## **Property      Meaning**

**Consistency** The database moves from one valid state to another, maintaining integrity constraints.

**Isolation** Concurrent transactions do not interfere with each other.

**Durability** Once a transaction is committed, it remains so even in the event of a failure.

◆ **Example:**

A banking transaction that transfers ₹100 from Account A to B:

- **Atomicity:** Both debit and credit occur, or neither.
  - **Consistency:** Balances are correctly updated.
  - **Isolation:** Other operations do not interfere.
  - **Durability:** Changes persist even if the system crashes after the commit.
- 

### **3. Eventual Consistency (Common in NoSQL Systems)**

◆ **Definition:**

Eventual consistency means that, **given enough time, all nodes in a distributed system will converge to the same data value**, assuming no new updates are made.

◆ **Characteristics:**

- Prioritizes **availability** over immediate consistency.
- Used in systems with **high scalability** needs and **loose latency requirements**.

◆ **Real-World Example:**

- **Amazon DynamoDB and Cassandra:**
  - A write may appear immediately on one node but take time to propagate to others.
  - Eventually, all replicas will agree.

◆ **Benefits:**

- **Low latency** responses.
- **High availability** during network partitions.
- **Scalability** across multiple regions.

◆ **Drawback:**

- Risk of **temporary stale reads** or conflicts if multiple writes happen at the same time.
-

**Comparison Table:**

Aspect	ACID	CAP Theorem	Eventual Consistency
Focus	Transaction integrity	Distributed system limitations	Long-term consistency of replicas
Trade-offs	Performance and scalability	Must choose 2 of 3 (C, A, P)	Sacrifices consistency temporarily
Used In	Relational DBs (SQL)	Distributed databases & systems	NoSQL DBs (DynamoDB, Cassandra)
Suitability	Banking, Finance, ERP	All large-scale systems	Social media, E-commerce, IoT